



**BITS Pilani**  
K K Birla Goa Campus



# 2. Boolean Retrieval

Information Retrieval

# Outline

---

## ❖ Evaluation of IR System

### Content

1. Term Vocabulary
2. Inverted Index
3. Boolean Retrieval
4. Faster Posting Lists
5. Extended Boolean Retrieval
6. Phrase Queries and Positional Postings

# Evaluation of an IR System



# Evaluation Metrics



	Retrieved	Not Retrieved
Relevant	Relevant Retrieved (tp)	Relevant Rejected (fn)
Not Relevant	Irrelevant Retrieved (fp)	Irrelevant Rejected (tn)


$$\textbf{Precision} = \frac{\textit{Relevant Retrieved}}{\textit{Retrieved}}$$

$$\textbf{Recall} = \frac{\textit{Relevant Retrieved}}{\textit{Relevant}}$$

# How good are the retrieved docs?

---

- **Precision** : Fraction of retrieved docs that are relevant to the user's **information need**
- **Recall** : Fraction of relevant docs in collection that are retrieved




**Question 1:** An IR system returns 8 relevant documents and 10 irrelevant documents. There are a total of 20 relevant documents in the corpus. What are the precision and recall on this search?

**Answer:**

Precision=  $8/18 = 0.44$

Recall=  $8/20 = 0.4$



**Question 2:** For a given query, there are 16 relevant documents in the collection. The precision of the query is 0.40 and the recall is 0.25. How many documents are there in the result set?

**Answer:**


$$\text{Recall} = \frac{\text{Relevant}}{16} = 0.25$$

$$\Rightarrow \text{Relevant} = 4$$

$$\text{Precision} = \frac{\text{Relevant [4]}}{\text{Retrieved}} = 0.4$$

$$\Rightarrow \text{Retrieved} = 10 \text{ (size of the result set)}$$

**Question 3:** Explain why both *precision* and *recall* metrics are required to evaluate an IR system and not just one metric individually?



**Collection:** 50 documents= 15 Relevant + 35 Irrelevant

- **High Recall and Low Precision:**

- All documents are retrieved irrespective of the query.

What we retrieved and what we missed.


- **High Precision and Low Recall:**

- Only 1 document (which is relevant) is retrieved irrespective of the query

- $Recall = \frac{\text{Relevant Retrieved}}{\text{Relevant}} = \frac{1}{15} = 0.06 = 6\%$

- $Precision = \frac{\text{Relevant Retrieved}}{\text{Retrieved}} = \frac{1}{1} = 1 = 100\%$





- Overall Precision
- Overall Recall
- Precision and Recall for top K documents

# F-measure

- Harmonic mean of Precision and Recall

- $Fscore = \frac{2PR}{P+R}$



R= 0.33

P= 0.67

$$Fscore = \frac{2 * 0.67 * 0.33}{0.67 + 0.33} = \frac{0.4422}{1} = 0.44$$

# 1. Term Vocabulary

---

- Tokenization
- Stopwords
- Normalization
- Lemmatization
- Stemming
- Thesaurus

# 1.1 Token and Terms

---

- Token is a particular instance of sequence of characters. Example:
    - Friends, and romans, and countrymen forming this one sentence.
    - Tokens:
      - Friends
      - And
      - Romans
      - And
      - Countrymen
      - Forming
      - this
      - One
      - Sentence
  - Term is a word in the dictionary.
    - Both tokens **and** will map with the same term **and** in the dictionary
    - What is being stored in the dictionary.
- Friend is a term
  - Friends is a token

# Issues in Tokenization

---

- **Finland's capital:**
  - Finland? Finlands? Finland's?
- **Hyphenated Sequence:**
  - Hewlett-Packard: Hewlett and Packard as two tokens?
  - State-of-the-art technology
  - Co-education
  - Lowercase, lower-case, lower case?
    - Query expansion
- **White spaces:**
  - San Francisco: one token or two?

# Issues in Tokenization

---

- **Numbers**

- 3/12/91
  - India: 3<sup>rd</sup> December
  - US: 12<sup>th</sup> March
  - Mar. 12, 1991
    - white space?
    - Other formats need to be preserved.
- B-52/ F-16- name of an aircraft or bomb
- SSN, PGP keys
- phone numbers
  - (832) 123-4567
  - Often have embedded spaces
  - Older IR systems may not index numbers
    - But often very useful- things like looking up error codes/stacktraces on the web
    - Can use n-grams
- Will often index *meta-data* separately
  - Creation date, format, etc.

# Issues in Tokenization

---

- **Language issues:**

- French

- L'ensemble: one token or two

- L? L'? Le?

- Want *l'ensemble* to match with *un ensemble*

- Until at least 2003, it didn't work on Google.

- German noun compounds are not segmented

- “Life insurance company employee”

- Lebensversicherungsgesellschaftsangestellter

- German retrieval systems benefit greatly from a compound splitter module

- Can give a 15% performance boost for German

# Issues in Tokenization

---

- **Language issues:**

- Chinese and Japanese have no spaces between words

汶川地震灾区首批自建永久性农房建成入住

- Not always guarantee a unique tokenization
- Further complicated in Japanese, with multiple alphabets intermingled

私はスポーツが大好きです

- Katakana, Hiragana, Kanji, romaji while end user can express query entirely in hiragana.



# Issues in Tokenization

---

- **Language issues:**

- In Korean language, insertion of spaces is flexible.

- Dates/궀 컴퓨터 게임
- January 17, 2018: 2018 년 1 월 17 일
- Number 17: seventeen: 십칠

# Issues in Tokenization



- **Language issues:**

- Arabic (or Hebrew) is written right to left but with certain items like numbers written left to right
- Words are separated but letters forms within a word form complex ligatures
- “Algeria achieved its independence in 1962 after 132 years of French occupation”
- حققت الجزائر استقلالها في عام 1962 بعد 132 عاما من الاحتلال الفرنسي
- With Unicode, the surface presentation is complex but the stored form is straightforward

# 1.2 Stop words

---

- With a stop list, you exclude from the dictionary entirely the **commonest words** that appear in almost all the documents.
  - They have little semantic content: the, a, an, and, to, be
  - There are a lot of them: ~30% of postings
- But the trend is away from doing this:
  - Good compression techniques means the space for including stop words in a system is very small
  - You need them for:
    - Phrase queries: “king of Denmark”
    - Various song titles: “let it be, to be or not to be”
    - Relational queries: “flights to London”

# 1.3 Normalization to Terms

---

- We need to "normalize" words in indexed text as well as query words into the same form
  - we want to match U.S.A and USA
  - Automation, automate, automatic
- Results is terms: a term is a (normalized) word type, which is an entry in our IR system dictionary
- We most commonly implicitly define equivalence classes of terms by, e.g.,
  - Deleting periods to form a term
    - U.S.A, USA { USA
  - Deleting hyphens to form a term
    - Anti-discriminatory, antidiscriminatory {antidiscriminatory

# Normalization: other languages



- Accents
  - French: résumé vs resume
  - Cliché vs cliche
- Umlauts
  - German: Tuebingen vs Tübingen
  - Should be equivalent
- Most important criterion:
  - How are your users like to write their queries for these words?
- Even in languages that standardly have accents, users often may not type them
  - Often best to normalize to a de-accented term
    - Tuebingen, Tübingen, Tübingen { Tübingen

# Normalization: other languages

---

- Normalization of things like date forms
  - 2018 년 1 월 18 일 vs 18/1
  - Japanese use of Kana vs Chinese characters
- Tokenization and normalization may depend on the language and so is intertwined with language detection
  - Morgen will ich in mit Is this German “mit” or MIT?
- Numbers: 7/10 or 10<sup>th</sup> July
- Crucial: need to normalize indexed text as well as the query terms into the same form

# Case Folding



- Reduce all letters to lower case
  - Query=automobile. A doc has a sentence starting with Automobile should be retrieved.
  - Exception: upper case in mid-sentence?
    - E.g.; General Motors
    - Fed vs fed
    - SAIL vs sail
  - Often best to lower case everything, since users will use lowercase regardless of ‘correct’ capitalization
- Google example:
  - Query C.A.T
  - First few results were for CAT exam and “cat” (animal) not Caterpillar Inc.

# Normalization to terms



- An alternative to equivalence classing is to do an asymmetric expansion (query expansion)
- An example of where it may be useful
  - Enter: **window** Search: **window, windows**
  - Enter: **windows** Search: **Windows, window, windows**
  - Enter: **Windows** Search: **Windows**
- Potentially more powerful, but less efficient



# 1.4 Lemmatization


---

- A sophisticated NLP technique for normalization
- In English language, words like "am, are, is" are formed from the base word (root or stem form) called "be"
  - car, cars, car's, cars' -> car
  - am, are, is, was -> be
- Reduce inflectional/variant forms to base form
  - the boy's cars are different colors-> the boy car be different color
- Lemmatization implies doing "proper" reduction to dictionary headword form

# 1.5 Stemming

---

- Lemmatization requires a fairly sophisticated linguistic analysis of the text
- Stemming is an alternative of lemmatization
- Reduce terms to their "root" before indexing
- Stemming suggests crude affix chopping
  - language dependent
  - e.g. automate(s), automatic, automation all reduced to automat.
  - The final terms may not always be a valid dictionary word



# Recap



- Precision
- Recall
- F-Score
- Term Vocabulary
  - Tokenization
  - Stopwords
  - Normalization
  - Lemmatization
  - Stemming

# Porter's Stemming



- commonest algorithm for stemming English
  - Results suggest it's at least as good as other stemming options
- Conventions +5 phases of reductions
  - phases applied sequentially
  - each phase consists of a set of commands
  - sample convention: of the rules in a compound command, select the one that applies to the longest suffix.
  - Goes through 5 phases of transformation before final term is generated

# Phase Example



- phase: a sequence of production rules
  - replacing a string appearing on left side with string appearing on right side
- has a bunch of rules in each phase
  - check if any of these rules has suffice in token
    - sses -> ss
    - ies -> i
    - ational -> ate
    - tional -> tion
- rules sensitive to the measure of words
  - (m>1) EMENT -> (m>1)
    - replacement -> replac
    - cement -> cement

# Other Stemmer



- Other stemmer exist, e.g. Lovins stemmer
  - <http://snowball.tartarus.org/algorithms/lovins/stemmer.html>
  - Single-pass, longest suffix removal (about 250 rules)
- Do stemming and other normalization help?
  - English: very mixed results. Helps recall but harms precision
    - operative (dentistry) -> oper
    - operational (research) -> oper
    - operating (systems) -> oper

# Language Specificity

---

- many of the above features embody transformations that are
  - language specific and
  - often application specific
- these are plug-in addenda to the indexing
- both open source and commercial plug-ins are available for handling these

# 1.6 Thesaurus

---

- Do we handle synonyms and homonyms?
  - e.g. by hand-constructed **equivalence classes**
    - **car=automobile**                      **color=colour**
  - we can rewrite to form equivalence-class terms
    - when the document contains automobile, index it under **car-automobile** (and vice-versa)
  - or we can use **query expansion**
    - when the query contains **automobile**, look under **car** as well



# Challenges



- Parsing a document
  - Format? Language? Character set? Unit document?
  - Complete book is a one document?
  - Each section or paragraph is a document?
  - Every sentence is a separate document?

## 2. Inverted Index




# Unstructured data in 1620



- Which plays of Shakespeare contain the words ***Brutus*** AND ***Caesar*** but ***NOT Calpurnia***?
- One could **grep** all of Shakespeare's plays for ***Brutus*** and ***Caesar***, then strip out lines containing ***Calpurnia***?
- Why is that not the answer?
  - Slow (for large corpora)
  - ***NOT Calpurnia*** is non-trivial
  - Other operations (e.g., find the word ***Romans*** near ***countrymen***) not feasible
  - Ranked retrieval (best documents to return)
    - Later lectures

# Term-document incidence matrices



	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

*Brutus AND Caesar BUT NOT  
Calpurnia*



1 if play contains word,  
0 otherwise

# Incidence vectors

- So we have a 0/1 vector for each term.
- To answer query: take the vectors for **Brutus**, **Caesar** and **Calpurnia** (complemented) → bitwise *AND*.
  - 110100 *AND* 110111 *AND* 101111 =
  - **100100**

**Limitations?**

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
<b>Antony</b>	1	1	1	1	1	1
<b>Brutus</b>	1	1	1	1	1	1
<b>Caesar</b>	1	1	1	1	1	1
<b>Calpurnia</b>	0	0	0	0	0	0
<b>Cleopatra</b>	1	1	1	1	1	1
<b>mercy</b>	1	1	1	1	1	1
<b>worser</b>	1	1	1	1	1	1

# Bigger collections

---

- Consider  $N = 1$  million documents, each with about 1000 words ( $10^3$  words).
- Average 6 bytes/word including spaces and punctuation
  - 6GB of data in the documents (1GB=  $10^9$  bytes).
- Say there are  $M = 500K$  *distinct* terms among these.

# Can't build the matrix



- 500K x 1M matrix ( $0.5 * 10^{12}$ ) has half-a-trillion 0's and 1's.
- But it has no more than one billion 1's.
  - matrix is extremely sparse.

# Inverted Index




- By far, the most commonly and frequently used data structure in modern IR systems.
- **INVERT** index?
- Earlier, a 2D metrics of terms that are present in a document- **sparsity**
- Now, for each term  $t$ , store a list of documents that contain  $t$ .



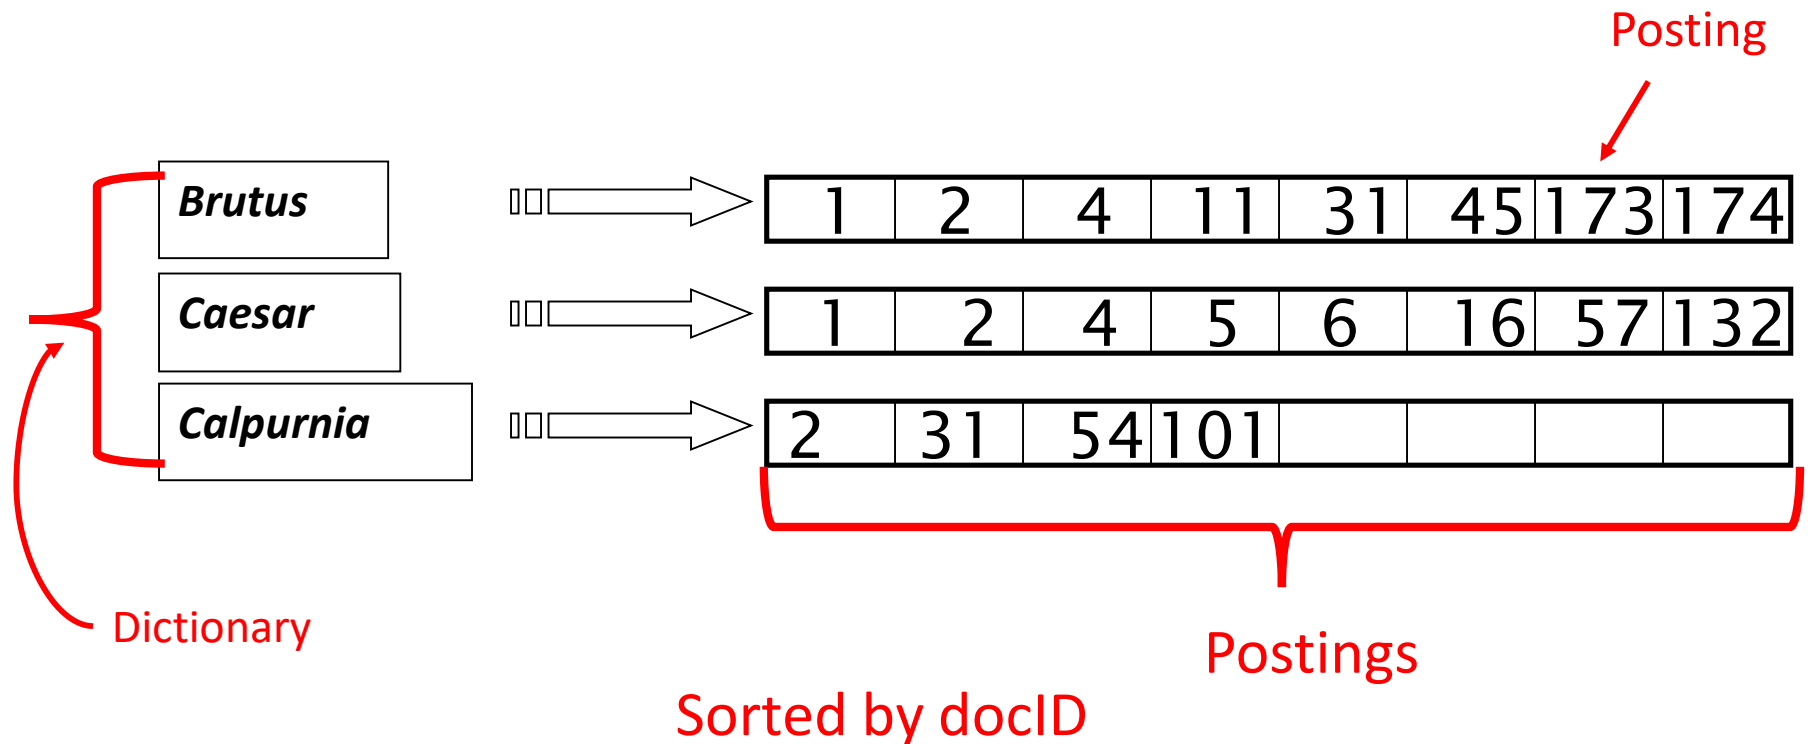
# Inverted index

- Assign a unique id to each document. Say, Doc1, Doc2, Doc3
- Sort the documents in a specific (say, ascending) fashion.



# Inverted index

- We need variable-size **postings lists**



# Inverted index construction

Documents to be indexed



Friends, Romans, countrymen.

⋮

Tokenizer

Token stream

Friends

Romans

Countrymen

Linguistic modules

Modified tokens

friend

roman

countryman


Indexer

Inverted index

*friend*

*roman*

*countryman*



# Indexer steps: Token sequence

- Sequence of (Modified token, Document ID) pairs.

Doc 1

I did enact Julius  
Caesar I was killed  
i' the Capitol;  
Brutus killed me.

Doc 2

So let it be with  
Caesar. The noble  
Brutus hath told you  
Caesar was ambitious



Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2
	44

# Indexer steps: Sort

- Sort by terms
  - And then docID

Term	docID		Term	docID
I	1		ambitious	2
did	1		be	2
enact	1		brutus	1
julius	1		brutus	2
caesar	1		capitol	1
I	1		caesar	1
was	1		caesar	2
killed	1		caesar	2
i'	1		did	1
the	1		enact	1
capitol	1		hath	1
brutus	1		I	1
killed	1		I	1
me	1		i'	1
so	2		it	2
let	2		julius	1
it	2		killed	1
be	2		killed	1
with	2		let	2
caesar	2		me	1
the	2		noble	2
noble	2		so	2
brutus	2		the	1
hath	2		the	2
told	2		told	2
you	2		you	2
caesar	2		was	1
was	2		was	2
ambitious	2		with	2

# Indexer steps: Dictionary & Postings

- Multiple term entries in a single document are merged.
- Split into Dictionary and Postings
- Doc. frequency information is added.

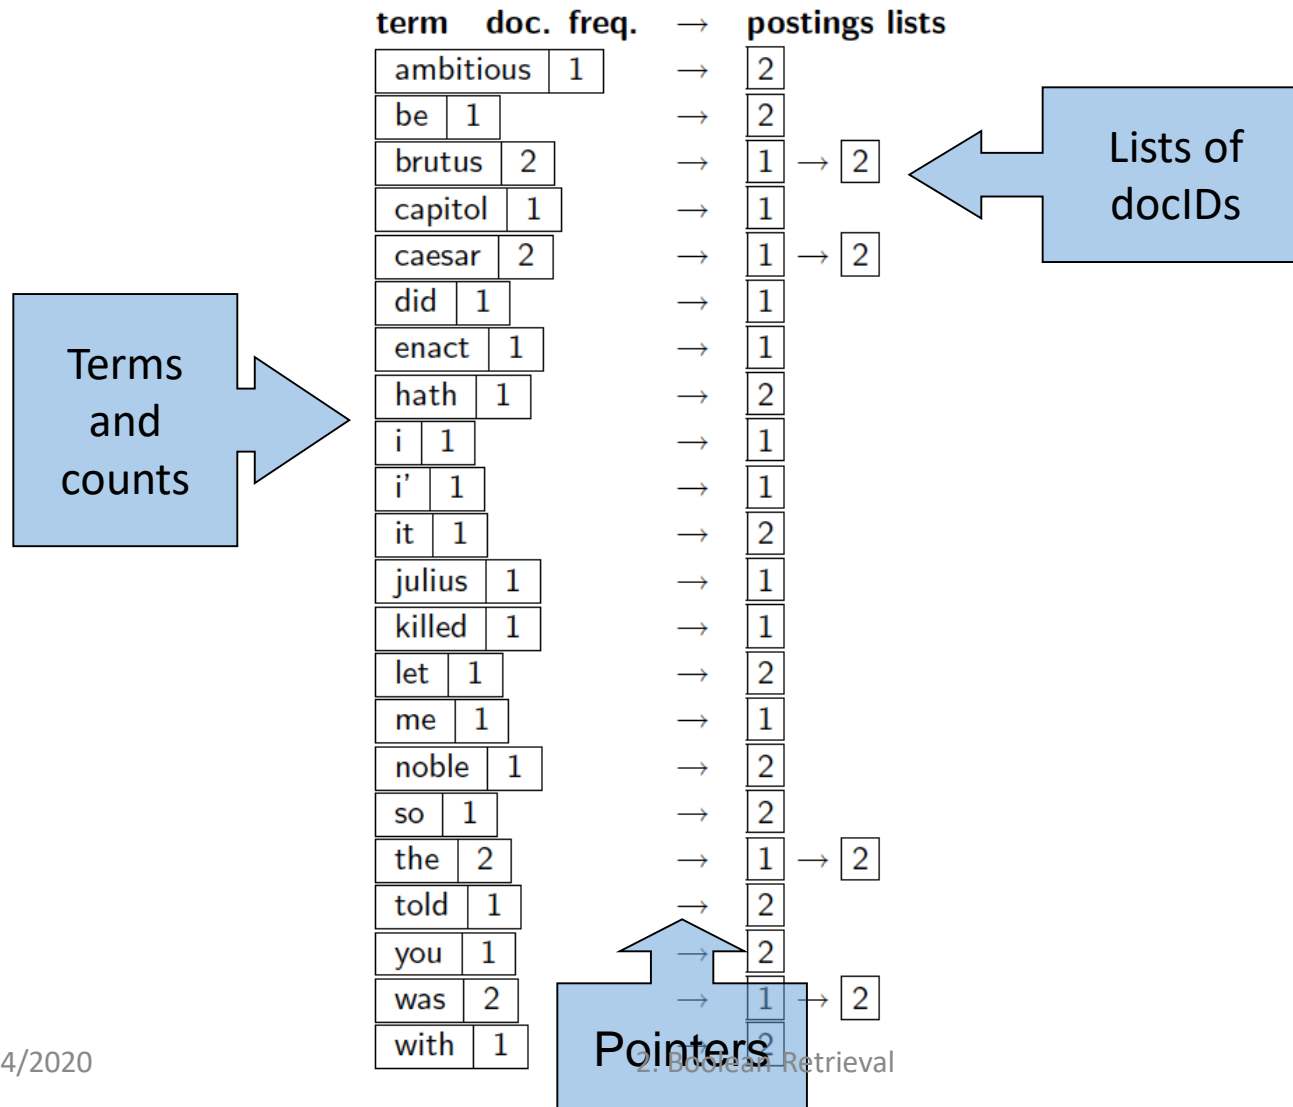
Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2



term	doc. freq.	→	postings lists
ambitious	1	→	[2]
be	1	→	[2]
brutus	2	→	[1] → [2]
capitol	1	→	[1]
caesar	2	→	[1] → [2]
did	1	→	[1]
enact	1	→	[1]
hath	1	→	[2]
i	1	→	[1]
i'	1	→	[1]
it	1	→	[2]
julius	1	→	[1]
killed	1	→	[1]
let	1	→	[2]
me	1	→	[1]
noble	1	→	[2]
so	1	→	[2]
the	2	→	[1] → [2]
told	1	→	[2]
you	1	→	[2]
was	2	→	[1] → [2]
with	1	→	[2]

Why frequency?  
Will discuss later.

# Where do we pay in storage?



# 3. The Boolean Retrieval Model





# Boolean Model

---

- Simple model based on **set theory and Boolean algebra**
  - Documents are sets of terms
  - Queries are Boolean expressions on terms
- Historically the most common model
- Queries specified as boolean expressions
  - Precise semantics
  - Neat formalism
- Terms are **either present or absent**. Thus,  $w_{ij} \in \{1,0\}$
- There are three connectives used: ***and, or, not***

# Boolean Model

---

- **D: set of words (indexing terms) present in a document**
  - each term is either present (1) or absent (0)
- **Q: A Boolean expression**
  - terms are index terms
  - operators are AND, OR, and NOT
- **F: Boolean algebra over sets of terms and sets of documents**
- **R: a document is predicted as relevant to a query expression if it *satisfies the query expression***
  - Example:  **$((\text{text} \vee \text{information}) \wedge \text{retrieval} \wedge \neg \text{theory})$**
- Each query term specifies a set of documents containing the term
- AND ( $\wedge$ ): the intersection of two sets
- OR ( $\vee$ ): the union of two sets
- NOT ( $\neg$ ): set inverse, or really set difference

# Boolean Queries: **Exact match**

---



- The **Boolean retrieval model** is being able to ask a query that is a Boolean expression:
  - Boolean Queries are queries using *AND*, *OR* and *NOT* to join query terms
    - Views each document as a set of words
    - Is precise: document matches condition or not.
  - Perhaps the simplest model to build an IR system on
- Primary commercial retrieval tool for 3 decades.
- Many search systems you still use are Boolean:
  - Email, library catalog, Mac OS X Spotlight

# Query processing: $x \text{ AND } y$

- Consider processing the query:

## ***Brutus AND Caesar***

- Locate ***Brutus*** in the Dictionary; Retrieve its postings.
- Locate ***Caesar*** in the Dictionary; Retrieve its postings.
- “Merge” the two postings (intersect the document sets):



**Documents to be retrieved.**

# Intersecting two postings lists (a “merge” algorithm)

```
INTERSECT( $p_1, p_2$ )  
1   $answer \leftarrow \langle \rangle$   
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$   
3  do if  $docID(p_1) = docID(p_2)$   
4      then  $\text{ADD}(answer, docID(p_1))$   
5           $p_1 \leftarrow next(p_1)$   
6           $p_2 \leftarrow next(p_2)$   
7      else if  $docID(p_1) < docID(p_2)$   
8          then  $p_1 \leftarrow next(p_1)$   
9          else  $p_2 \leftarrow next(p_2)$   
10 return  $answer$ 
```

# Query Processing: $x$ OR $y$

Algorithm for AND query.

INTERSECT( $p_1, p_2$ )

Changes for OR query.

```
1  answer  $\leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then ADD(answer,  $\text{docID}(p_1)$ )
5           $p_1 \leftarrow \text{next}(p_1)$ 
6           $p_2 \leftarrow \text{next}(p_2)$ 
```

```
Else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
    then add (answer,  $\text{docID}(p_1)$ );     $p_1 \leftarrow \text{next}(p_1)$ ;
Else
    then add (answer,  $\text{docID}(p_2)$ );     $p_2 \leftarrow \text{next}(p_2)$ ;
```

```
10 return answer
```


# READING TASK- 1

---

- Query Processing algorithm:
  - $x \text{ AND } (\text{NOT } y)$

# Query Optimization

- What is the best order for query processing?
- Consider a query that is an *AND* of  $n$  terms.
- For each of the  $n$  terms, get its postings, then *AND* them together.



**Query: *Brutus AND Calpurnia AND Caesar***



# Query optimization example

- Process in order of increasing freq:
  - *start with smallest set, then keep cutting further.*

This is why we kept  
document freq. in dictionary

Brutus	⇒	2	4	8	16	32	64	128	
Caesar	⇒	1	2	3	5	8	16	21	34
Calpurnia	⇒	13	16						

Execute the query as (***Calpurnia AND Brutus***) ***AND Caesar***.

# Exercise- OR query

- Recommend a query processing order for

*(tangerine OR trees) AND  
(marmalade OR skies) AND  
(kaleidoscope OR eyes)*

- Which two terms should we process first?

Term	Freq
eyes	213312
kaleidoscope	87009
marmalade	107913
skies	271658
tangerine	46653
trees	316812

# Query processing exercises

---

- **Exercise:** If the query is *friends AND romans AND (NOT countrymen)*, how could we use the freq of *countrymen*?
- **Exercise:** Extend the merge to an arbitrary Boolean query. Can we always guarantee execution in time linear in the total postings size?
- **Hint:** Begin with the case of a Boolean *formula* query: in this, each query term appears only once in the query.

# Advantages



- Clean Formalism
- Easy to implement
- Intuitive concept
- Still it is dominant model for document database systems

# Limitations

---

- Retrieval based on binary decision criteria with no notion of partial matching
- No ranking of the documents is provided (absence of a grading scale)
- Information need has to be translated into a Boolean expression which most users find awkward
- The Boolean queries formulated by the users are most often too simplistic
- Frequently returns either too few or too many documents in response to a user query

# Exercise

---


- Try the search feature at <http://www.rhymezone.com/shakespeare/>
- Write down five search features you think it could do better

# 4. Faster Posting Lists



# Faster Posting Lists


- **Faster postings list intersection via skip pointers**
  - ways to increase the efficiency of using postings lists





# Skip Pointers

- Augmenting postings lists with skip pointers
- Skip pointers are effectively shortcuts that allow us to avoid processing parts of the postings list that will not figure in the search results.
- Where to place skip pointers ?
- How to do efficient merging using skip pointers?



# Postings lists intersection with skip pointers

IntersectWithSkips(p1, p2)

1 answer  $\leftarrow$   $\langle \rangle$

2 while p1  $\neq$  NIL and p2  $\neq$  NIL

3 do if docID(p1) = docID(p2) then

4         ADD(answer, docID(p1))

5                 p1  $\leftarrow$  next(p1)

6                 p2  $\leftarrow$  next(p2)

7         else if docID(p1) < docID(p2) then

8                 if hasSkip(p1) and (docID(skip(p1))  $\leq$  docID(p2)) then

9                         while hasSkip(p1) and (docID(skip(p1))  $\leq$  docID(p2))

10                                 do p1  $\leftarrow$  skip(p1)

11                         else p1  $\leftarrow$  next(p1)

12                 else if hasSkip(p2) and (docID(skip(p2))  $\leq$  docID(p1)) then

13                         while hasSkip(p2) and (docID(skip(p2))  $\leq$  docID(p1))

14                                 do p2  $\leftarrow$  skip(p2)

15                 else p2  $\leftarrow$  next(p2)

16 return answer

# Skip Pointers



- Available only for the original postings
- For an intermediate result in a complex query, the call `haskip(p)` will always return false
- The presence of skip pointers only helps for AND queries, not for OR queries.

# Placing Skips

---

- Simple heuristic: for postings of length  $P$ , use  $\lceil P \rceil$  evenly-spaced skip pointers.
- This ignores the distribution of query terms.
- Easy if the index is relatively static; harder if  $P$  keeps changing because of updates.
- Building effective skip pointers is easy if an index is relatively static
- Difficult if a postings list keeps changing because of updates

# 5. Phrase Queries and Positional Postings



# Phrase Queries

---

- Want to be able to answer queries such as “Stanford University” – as a phrase
  - Thus the sentence “I went to university at Stanford” is not a match.
  - The sentence “The inventor Stanford Ovshinsky never went to university” is not a match.
- The concept of phrase queries has proven easily understood by users; one of the few “advanced search” ideas that works
- Many more queries are implicit phrase queries

# Handling Phrase Queries



- For this, it no longer suffices to store only
  - <term : docs> entries
- The two approaches to support phrase queries and their combination are
  1. Biword Indexes
  2. Positional Indexes
- A search engine should not only support phrase queries, but implement them efficiently.

# 5.1 Biword Indexes

---

- Index every consecutive pair of terms in the text as a phrase
- For example the text “Friends, Romans, Countrymen” would generate the biwords:
  - *friends romans*
  - *romans countrymen*
- Each of these biwords is now a dictionary term
- Two-word phrase query-processing is now immediate.
- Longer phrases are processed by breaking them.
- “*stanford university palo alto*” can be broken into the Boolean query on biwords:
  - *stanford university AND university palo AND palo alto*
- Without the docs, we cannot verify that the docs matching the above Boolean query do contain the phrase.



# Extended Biwords



- Parse the indexed text and perform part-of-speech-tagging (POST).
- Bucket the terms into (say) Nouns (N) and articles/prepositions (X).
- Call any string of terms of the form  $NX^*N$  an extended biword.
  - Each such extended biword is now made a term in the dictionary.

# Example



- cost overruns on a power plant

N      N      X      X      N      N

- Query processing: parse it into N's and X's
- Segment query into enhanced biwords
- “cost overruns” AND “overruns power” AND “power plant”
- Better results can be obtained by using more precise part-of-speech patterns that define which extended biwords should be indexed.

# Limitations



- False positives
- Index blowup due to bigger dictionary
- Biword indexes are not the standard solution (for all biwords) but can be part of a compound strategy.

## 5.2 Positional Indexes

- In the postings, store, for each ***term*** the position(s) in which tokens of it appear:
  - <***term***, number of docs containing ***term***;
  - *doc1*: frequency of the term; position1, position2 ... ;
  - *doc2*: frequency of the term; position1, position2 ... ;
  - etc.>

```
to, 993427: (  
1, 6: (7, 18, 33, 72, 86, 231);  
2, 5: (1, 17, 74, 222, 255);  
4, 5: (8, 16, 190, 429, 433);  
5, 2: (363, 367);  
7, 3: (13, 23, 191); ...  
...  
)
```

```
be, 178239: (  
1, 2: (17, 25);  
4, 5: (17, 191, 291, 430, 434);  
5, 3: (14, 19, 101); ...  
...  
...  
)
```

# Positional index

---

- For phrase queries, we use a merge algorithm recursively at the document level.
- This requires more than just equality.
- Start with the least frequent term and then work to further restrict the list of possible candidates.
- Need to check that their positions of appearance in the document are compatible with the phrase query being evaluated.
- This requires working out offsets between the words
- Positional Index can also be used for Proximity Searches

# Proximity Queries

---

- **Information need:** *information on the legal theories involved in preventing the disclosure of trade secrets by employees formerly employed by a competing company.*
- **Query:** "trade secret" /s disclos! /s prevent /s employe!
- **Information need:** *what are the requirements for disabled people to be able to access a workplace.*
- **Query:** disab! /p access! /s work-site work-place (employment /3 place)
- **LIMIT!** /3 STATUTE /3 FEDERAL /2 TORT
  - Again, here, /k means “within  $k$  words of”.
- Clearly, positional indexes can be used for such queries; biword indexes cannot.
- Exercise: Adapt the linear merge of postings to handle proximity queries. Can you make it work for any value of  $k$ ?
  - This is a little tricky to do correctly and efficiently